

# NY flights

April 16, 2023

```
[2]: library(nycflights13)
library(ggplot2)
library(tidyverse)

# 1.1)
# create a new column in "flights" with the date of departure
flights$date <- as.Date(with(flights, paste(day, month, year, sep="-")),
  ↪ "%d-%m-%Y")
```

Attaching core tidyverse packages tidyverse  
2.0.0

dplyr	1.1.0	readr	2.1.4
forcats	1.0.0	stringr	1.5.0
lubridate	1.9.2	tibble	3.2.0
purrr	1.0.1	tidyr	1.3.0

Conflicts

```
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag() masks stats::lag()
Use the conflicted package
(<http://conflicted.r-lib.org/>) to force all conflicts to
become errors
```

```
[3]: # create a new data.frame with 3 columns: origin airport, date and number of
  ↪ flights
from_ny <- flights %>% count(origin, date)


head(from_ny)
```

A tibble: 6 × 3

origin	date	n
<chr>	<date>	<int>
EWR	2013-01-01	305
EWR	2013-01-02	350
EWR	2013-01-03	336
EWR	2013-01-04	339
EWR	2013-01-05	238
EWR	2013-01-06	301

```
[4]: # plot of the total number of flights departed from each of the three NYC
      ↪airports as a function of time
ggp_from_ny <- ggplot(from_ny, aes(date, n, col = origin)) + geom_point() +
      ↪theme_bw() + theme(legend.position = "top") + theme(legend.title =
      ↪element_text(size=14, face = "bold")) + ggtitle("Number of flights") +
      ↪labs(y = "n", x = "date")

ggp_from_ny
```



```
[5]: # 1.2)
      # create a function that gives the number of the week from the date
count_week <- function(x){
  n <- rep(0, length(x))
```

```

for (i in 1:length(x)) {
  if (x[i] == "2013-12-30" || x[i] == "2013-12-31"){
    n[i] <- 53
  }
  else {
    n[i] <- isoweek(x[i])
  }
}
return(n)
}

# create 2 new columns in "flights" with the number of the day and the number
  ↳ of the week
flights$weekday <- wday(flights$date, week_start = 1)
flights$week_num <- count_week(flights$date)

```

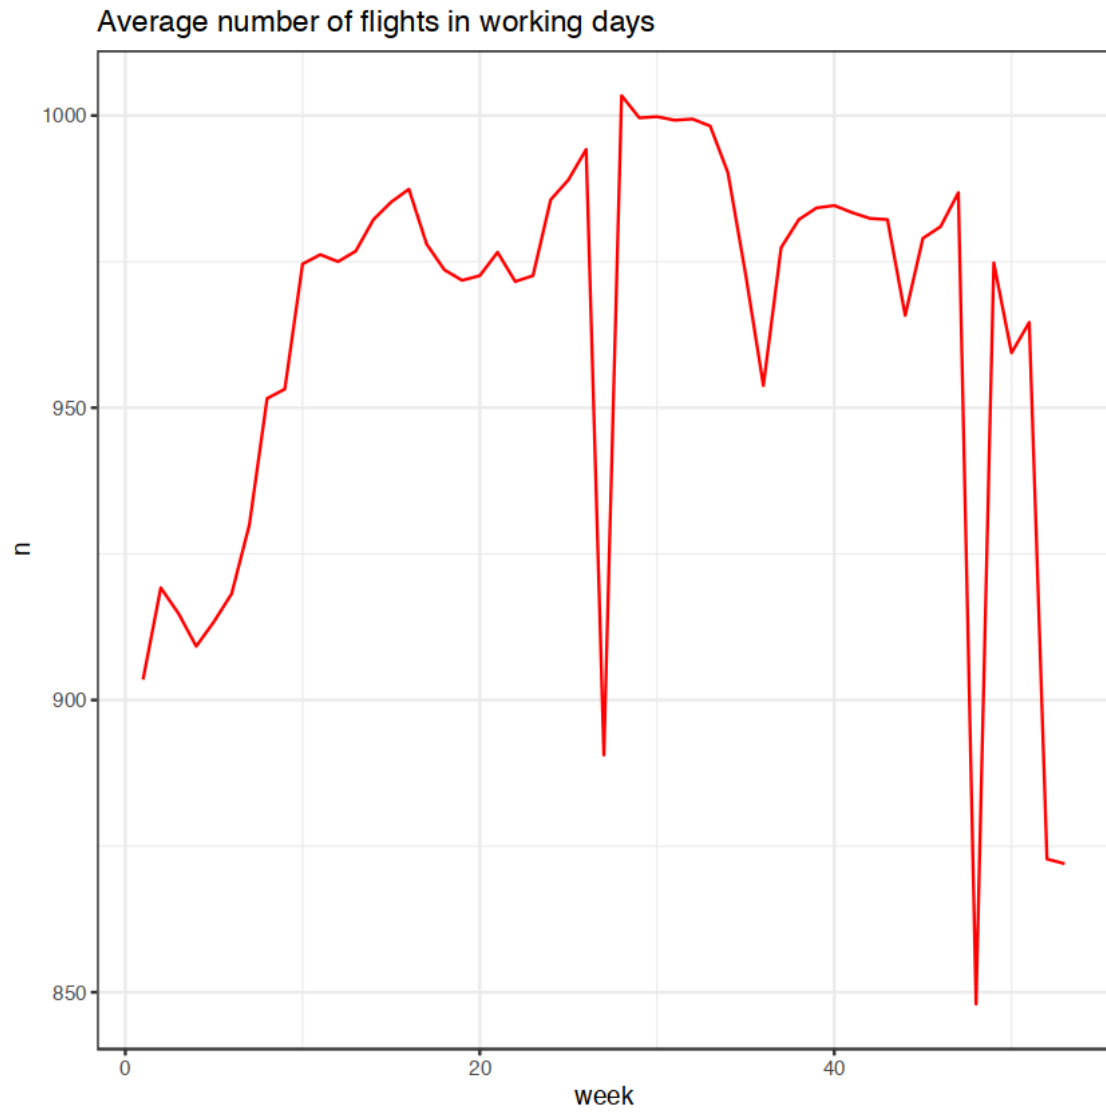
```

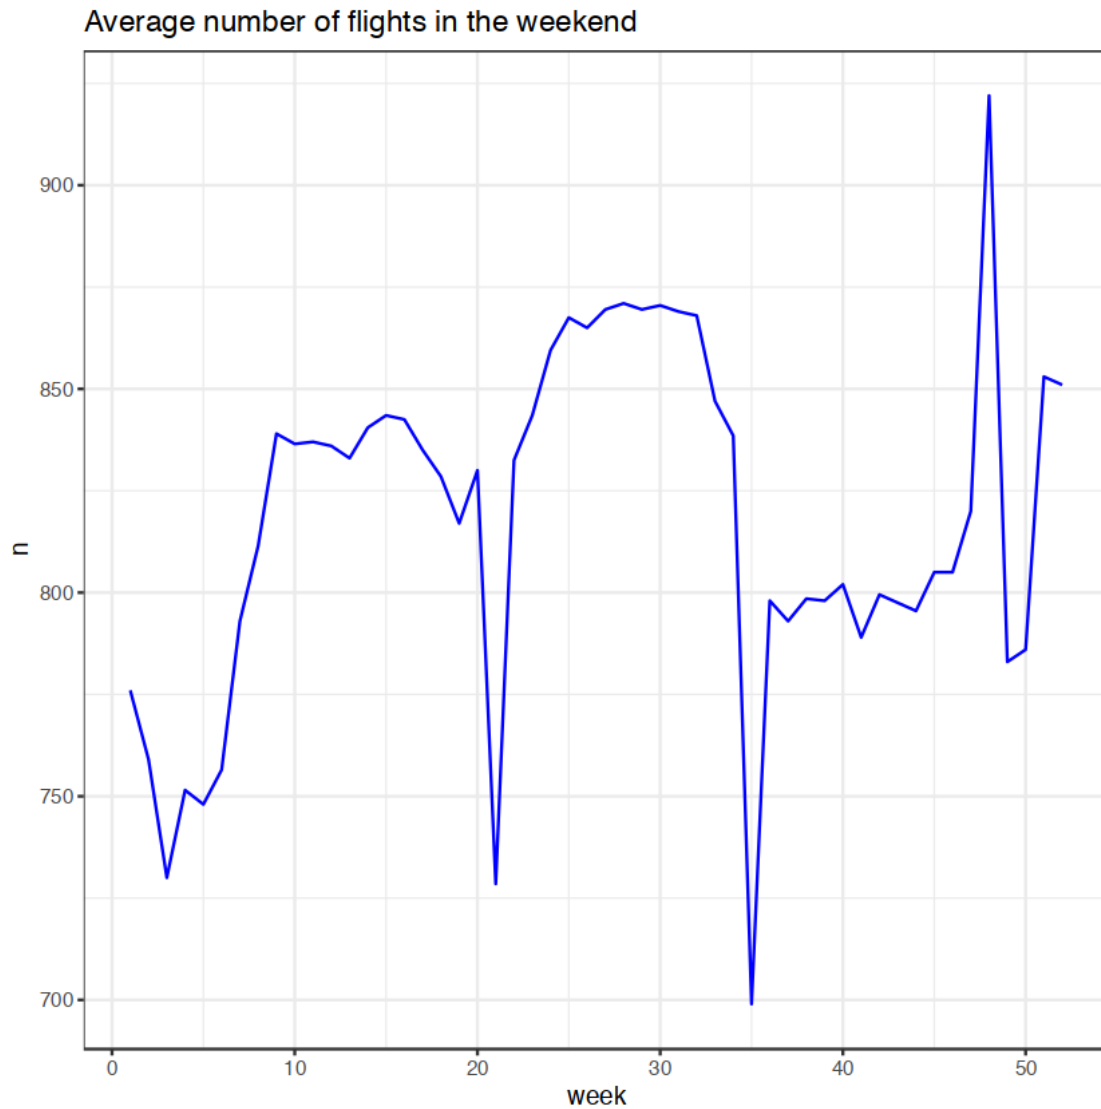
[6]: #divide the working days from the weekend
working_day <- filter(flights, weekday %in% c(1,2,3,4,5))
weekend_day <- filter(flights, weekday %in% c(6,7))

# calculate the mean number of flights for every week
working_day <- working_day %>% count(week_num, weekday) %>% group_by(week_num)
  ↳ %>% summarise(mean=mean(n), .groups = 'drop')
weekend_day <- weekend_day %>% count(week_num, weekday) %>% group_by(week_num)
  ↳ %>% summarise(mean=mean(n), .groups = 'drop')

# create plot
ggp_average_working <- ggplot(working_day, aes(week_num, mean)) +
  ↳ geom_line(color = "red") + theme_bw() + ggtitle("Average number of flights
  ↳ in working days") + labs(y = "n", x = "week")
ggp_average_weekend <- ggplot(weekend_day, aes(week_num, mean)) +
  ↳ geom_line(color = "blue") + theme_bw() + ggtitle("Average number of flights
  ↳ in the weekend") + labs(y = "n", x = "week")
ggp_average_working
ggp_average_weekend

```





```
[46]: # 2.1)
# create 3 new tibbles with 3 columns: date, origin airport and min/max/average
# delay
min_delay <- flights %>% group_by(date, origin) %>% summarise(min =
# dep_delay[which.min(dep_delay)], .groups = 'drop')
max_delay <- flights %>% group_by(date, origin) %>% summarise(max =
# dep_delay[which.max(dep_delay)], .groups = 'drop')
average_delay <- flights %>% group_by(date, origin) %>% summarise(average =
# mean(dep_delay, na.rm = TRUE), .groups = 'drop')


# create plot
```

```


ggp_min_delay <- ggplot(min_delay, aes(date, min, col = origin)) + geom_line()
  ↳+ theme_bw() + theme(legend.position = "top") + theme(legend.title =
  ↳element_text(size=14, face = "bold")) + ggtitle("Minimum delay") + labs(y =
  ↳"delay", x = "date")
ggp_max_delay <- ggplot(max_delay, aes(date, max, col = origin)) + geom_line()
  ↳+ theme_bw() + theme(legend.position="top") + theme(legend.title =
  ↳element_text(size=14, face = "bold")) + ggtitle("Maximum delay") + labs(y =
  ↳"delay", x = "date")
ggp_average_delay <- ggplot(average_delay, aes(date, average, col = origin)) +
  ↳geom_line() + theme_bw() + theme(legend.position="top") + theme(legend.title
  ↳= element_text(size = 14, face = "bold")) + ggtitle("Average delay") +
  ↳labs(y = "delay", x = "date")

ggp_min_delay
ggp_max_delay
ggp_average_delay

```




## Maximum delay



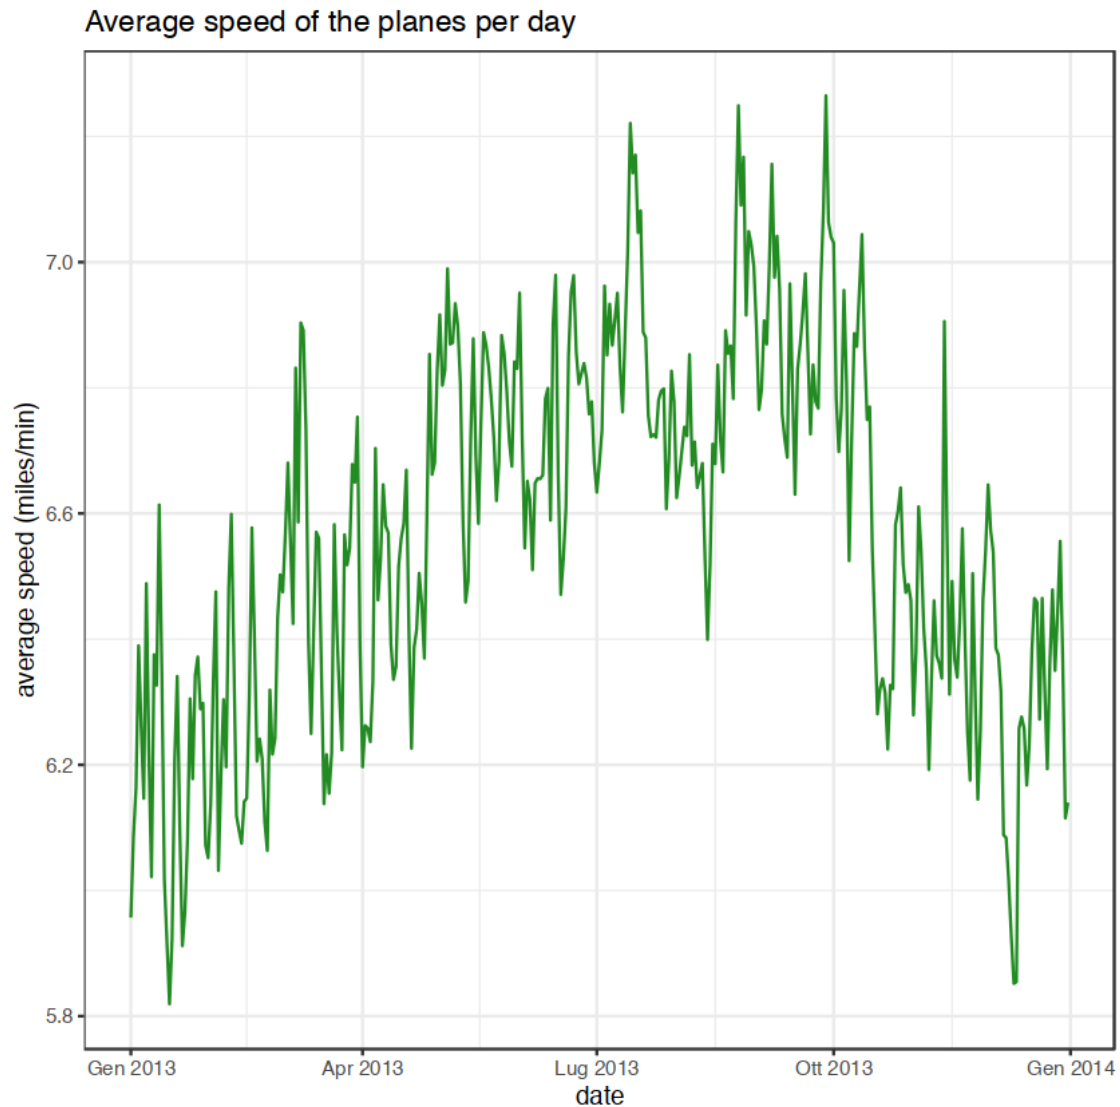


## Average delay



```
[10]: # 3)
# add a new column with the average speed (miles/min) of every flight and
# calculate the mean for every day
flights$speed <- flights$distance / flights$air_time
speed <- flights %>% group_by(date) %>% summarise(average = mean(speed, na.rm =
# TRUE), .groups = 'drop')

# create plot
ggp_speed <- ggplot(speed, aes(date, average)) + geom_line(color = "forest_
# green") + theme_bw() + ggtitle("Average speed of the planes per day") +
# labs(y = "average speed (miles/min)", x = "date")
ggp_speed
```



```
[12]: # 4.1)
# create two tibbles with date/number of the week, company and number of
# flights of every company per day/week
company_day <- subset(flights, select = c(carrier, date)) %>% count(carrier,
# date)
company_day <- company_day[with(company_day, order(date, n)),]

company_week <- subset(flights, select = c(carrier, week_num)) %>%
# count(carrier, week_num)
company_week <- company_week[with(company_week, order(week_num, n)),]

head(company_day)
head(company_week)
```

A tibble: 6 × 3	carrier	date	n
	<chr>	<date>	<int>
	HA	2013-01-01	1
	AS	2013-01-01	2
	F9	2013-01-01	2
	FL	2013-01-01	10
	VX	2013-01-01	12
	WN	2013-01-01	27

A tibble: 6 × 3	carrier	week_num	n
	<chr>	<dbl>	<int>
	YV	1	5
	HA	1	6
	AS	1	12
	F9	1	12
	FL	1	62
	VX	1	72

```
[13]: # create 2 new tibbles with 4 columns: date/number of the week, largest two
      ↪ numbers of flights and corresponding airports
first_two_per_day <- company_day %>% group_by(date) %>% summarise(max_2 =
      ↪ n[length(n) - 1], company_2 = carrier[length(n) - 1], max = n[length(n)],
      ↪ company= carrier[length(n)], .groups = 'drop')
first_two_per_week <- company_week %>% group_by(week_num) %>% summarise(max_2 =
      ↪ n[length(n) - 1], company_2 = carrier[length(n) - 1], max = n[length(n)],
      ↪ company= carrier[length(n)], .groups = 'drop')

head(first_two_per_day)
head(first_two_per_week)
```

A tibble: 6 × 5	date	max_2	company_2	max	company
	<date>	<int>	<chr>	<int>	<chr>
	2013-01-01	163	B6	165	UA
	2013-01-02	162	B6	170	UA
	2013-01-03	159	UA	162	B6
	2013-01-04	161	B6	161	UA
	2013-01-05	117	UA	154	B6
	2013-01-06	137	UA	156	B6
A tibble: 6 × 5	week_num	max_2	company_2	max	company
	<dbl>	<int>	<chr>	<int>	<chr>
	1	909	UA	958	B6
	2	994	B6	1035	UA
	3	970	B6	1032	UA
	4	960	B6	1032	UA
	5	963	B6	1039	UA
	6	972	B6	1041	UA

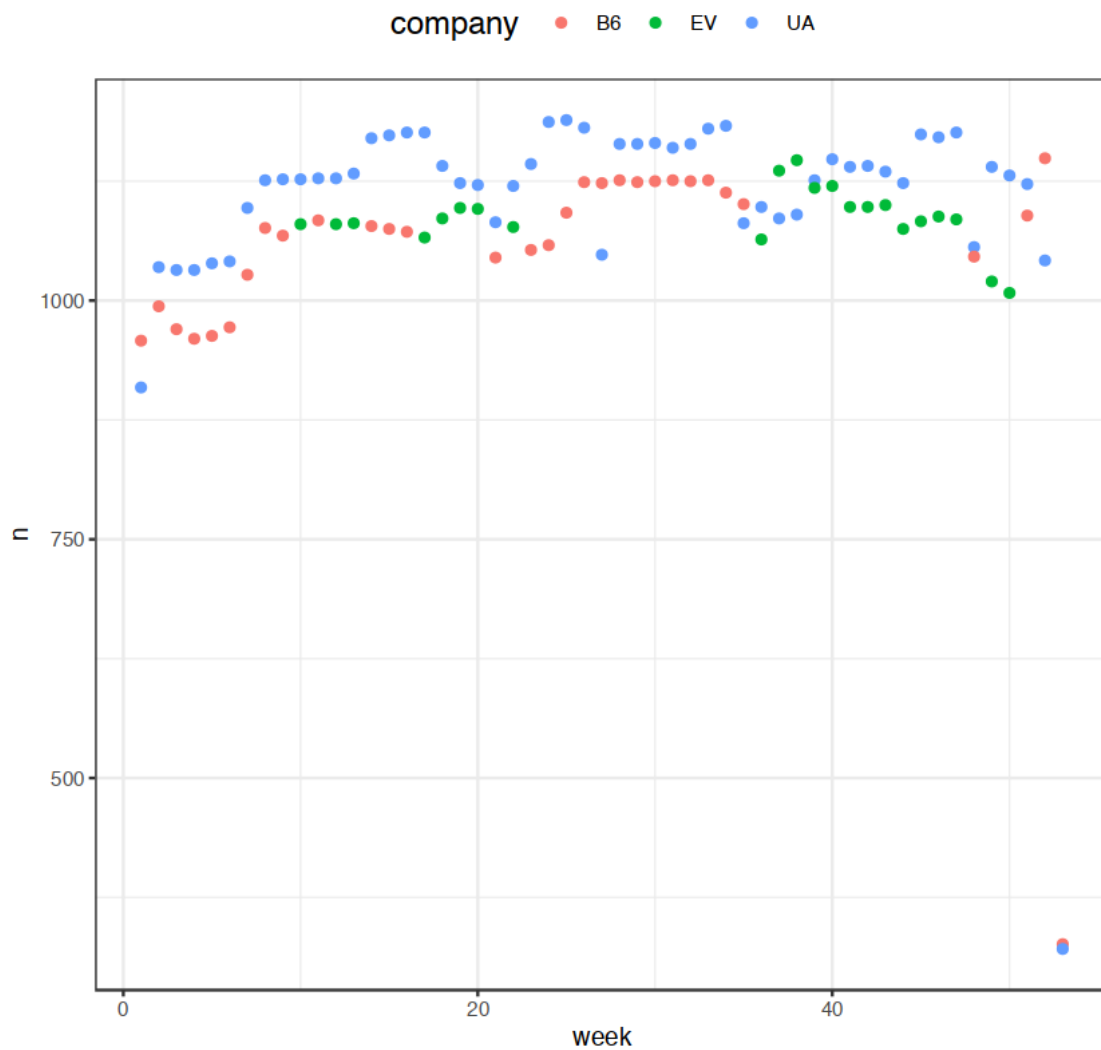
```
[14]: # new data frames to create a single plot
first_two_per_day <- data.frame(date = first_two_per_day$date,
                                max_two = c(first_two_per_day$max,
↪first_two_per_day$max_2),
                                company = c(first_two_per_day$company,
↪first_two_per_day$company_2))
first_two_per_week <- data.frame(week = first_two_per_week$week_num,
                                max_two = c(first_two_per_week$max,
↪first_two_per_week$max_2),
                                company = c(first_two_per_week$company,
↪first_two_per_week$company_2))

# create plot
ggp_first_two_per_day <- ggplot(first_two_per_day, aes(date, max_two, col =
↪company)) + geom_point() + theme_bw() + theme(legend.position="top") +
↪theme(legend.title = element_text(size=14, face = "bold")) +
↪ggtitle("Largest two numbers of flights per day") + labs(y = "n", x = "date")
ggp_first_two_per_week <- ggplot(first_two_per_week, aes(week, max_two, col =
↪company)) + geom_point() + theme_bw() + theme(legend.position="top") +
↪theme(legend.title = element_text(size=14, face = "bold")) +
↪ggtitle("Largest two numbers of flights per week") + labs(y = "n", x =
↪"week")
ggp_first_two_per_day
ggp_first_two_per_week
```

Largest two numbers of flights per day



Largest two numbers of flights per week



```
[15]: # 4.2)
# create a new data.frame with 3 columns: month, company and number of flights
company_month <- flights %>% count(month, carrier)

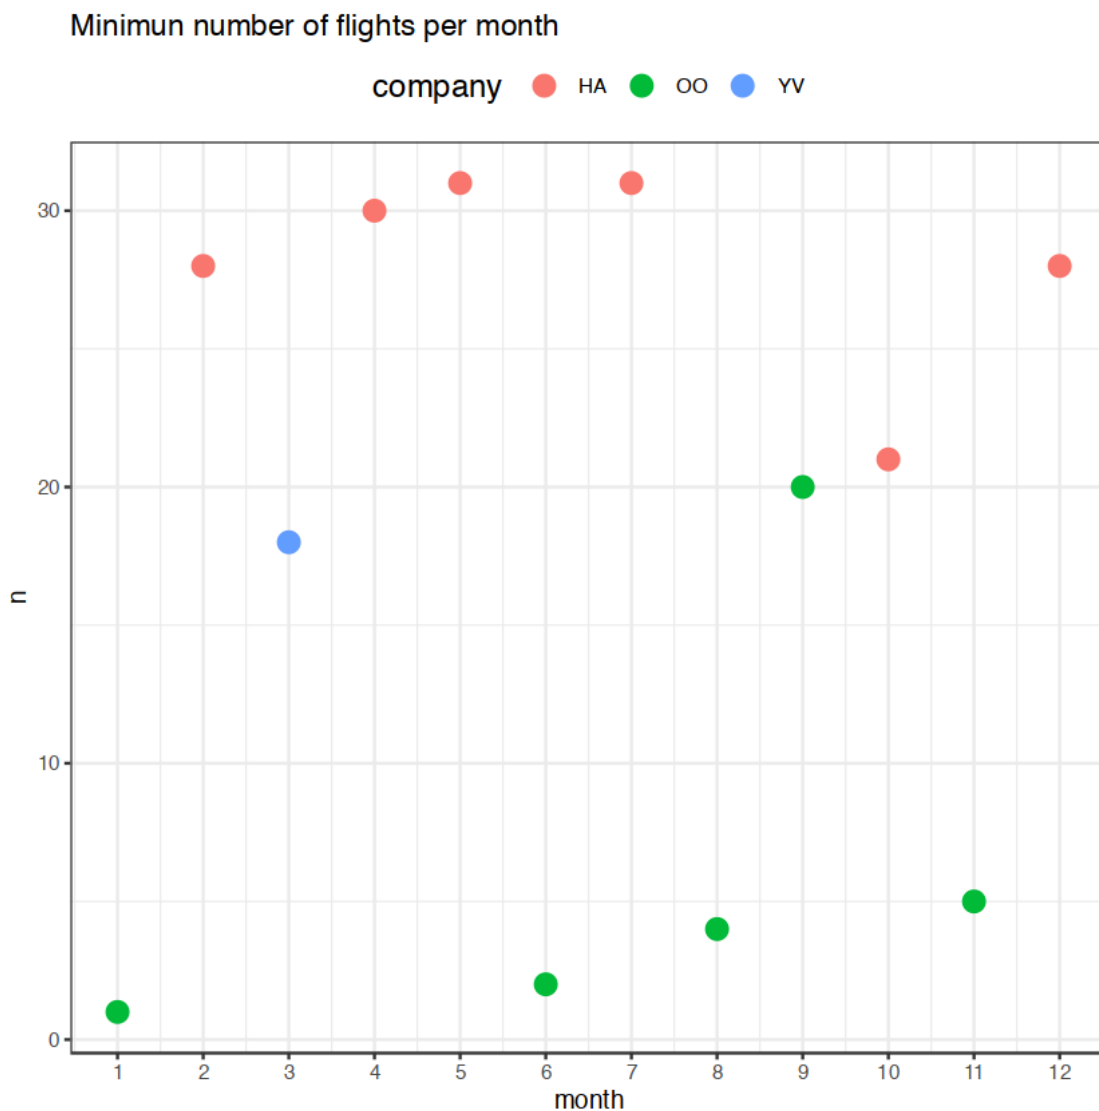
head(company_month)
```

month	carrier	n
<int>	<chr>	<int>
1	9E	1573
1	AA	2794
1	AS	62
1	B6	4427
1	DL	3690
1	EV	4171

A tibble: 6 × 3

```
[24]: # create a new tibble with 3 columns: month, company and minimum number of
      ↪ flights
min_per_month <- company_month %>% group_by(month) %>% summarise(min = n[which.
      ↪ min(n)], company = carrier[which.min(n)], .groups = 'drop')

# create plot
ggp_min_per_month <- ggplot(min_per_month, aes(month, min, col = company)) +
  ↪ geom_point(size = 4) + theme_bw() + theme(legend.position="top") +
  ↪ theme(legend.title = element_text(size=14, face = "bold")) +
  ↪ ggtitle("Minimun number of flights per month") + labs(y = "n", x = "month")
  ↪ + scale_x_continuous(breaks = scales::pretty_breaks(n = 12))
ggp_min_per_month
```



```
[48]: # 4.3)
# create a new tibble with month, maximum distance and company
distance_month <- subset(flights, select = c(carrier, month, distance))
distance_month <- distance_month %>% group_by(month) %>% summarise(max_distance =
  distance[which.max(distance)], company = carrier[which.max(distance)], .
  groups = 'drop')

distance_month
```

A tibble: 12 × 3

month	max_distance	company
<int>	<dbl>	<chr>
1	4983	HA
2	4983	HA
3	4983	HA
4	4983	HA
5	4983	HA
6	4983	HA
7	4983	HA
8	4983	HA
9	4983	HA
10	4983	HA
11	4983	HA
12	4983	HA